



```

In [0]: import collections
import cv2
import numpy as np
import matplotlib.pyplot as plt
import gym

def plot_learning_curve(x, scores, epsilons, filename, lines=None):
    fig=plt.figure()
    ax=fig.add_subplot(111, label="1")
    ax2=fig.add_subplot(111, label="2", frame_on=False)

    ax.plot(x, epsilons, color="C0")
    ax.set_xlabel("Training Steps", color="C0")
    ax.set_ylabel("Epsilon", color="C0")
    ax.tick_params(axis='x', colors="C0")
    ax.tick_params(axis='y', colors="C0")

    N = len(scores)
    running_avg = np.empty(N)
    for t in range(N):
        running_avg[t] = np.mean(scores[max(0, t-20):(t+1)])

    ax2.scatter(x, running_avg, color="C1")
    ax2.axes.get_xaxis().set_visible(False)
    ax2.yaxis.tick_right()
    ax2.set_ylabel('Score', color="C1")
    ax2.yaxis.set_label_position('right')
    ax2.tick_params(axis='y', colors="C1")

    if lines is not None:
        for line in lines:
            plt.axvline(x=line)

    plt.savefig(filename)

class RepeatActionAndMaxFrame(gym.Wrapper):
    def __init__(self, env=None, repeat=4, clip_reward=False, no_ops=0,
                 fire_first=False):
        super(RepeatActionAndMaxFrame, self).__init__(env)
        self.repeat = repeat
        self.shape = env.observation_space.low.shape
        self.frame_buffer = np.zeros_like((2, self.shape))
        self.clip_reward = clip_reward
        self.no_ops = no_ops
        self.fire_first = fire_first

    def step(self, action):
        t_reward = 0.0
        done = False
        for i in range(self.repeat):
            obs, reward, done, info = self.env.step(action)
            if self.clip_reward:
                reward = np.clip(np.array([reward]), -1, 1)[0]
            t_reward += reward
            idx = i % 2

```

```

        self.frame_buffer[idx] = obs
    if done:
        break

max_frame = np.maximum(self.frame_buffer[0], self.frame_buffer[1])
return max_frame, t_reward, done, info

def reset(self):
    obs = self.env.reset()
    no_ops = np.random.randint(self.no_ops)+1 if self.no_ops > 0 else 0
    for _ in range(no_ops):
        _, _, done, _ = self.env.step(0)
        if done:
            self.env.reset()
    if self.fire_first:
        assert self.env.unwrapped.get_action_meanings()[1] == 'FIRE'
        obs, _, _, _ = self.env.step(1)

    self.frame_buffer = np.zeros_like((2,self.shape))
    self.frame_buffer[0] = obs

    return obs

class PreprocessFrame(gym.ObservationWrapper):
    def __init__(self, shape, env=None):
        super(PreprocessFrame, self).__init__(env)
        self.shape = (shape[2], shape[0], shape[1])
        self.observation_space = gym.spaces.Box(low=0.0, high=1.0,
                                                shape=self.shape, dtype=np.float32)

    def observation(self, obs):
        new_frame = cv2.cvtColor(obs, cv2.COLOR_RGB2GRAY)
        resized_screen = cv2.resize(new_frame, self.shape[1:],
                                    interpolation=cv2.INTER_AREA)
        new_obs = np.array(resized_screen, dtype=np.uint8).reshape(self.shape)
        new_obs = new_obs / 255.0

        return new_obs

class StackFrames(gym.ObservationWrapper):
    def __init__(self, env, repeat):
        super(StackFrames, self).__init__(env)
        self.observation_space = gym.spaces.Box(
            env.observation_space.low.repeat(repeat, axis=0),
            env.observation_space.high.repeat(repeat, axis=0),
            dtype=np.float32)
        self.stack = collections.deque(maxlen=repeat)

    def reset(self):
        self.stack.clear()
        observation = self.env.reset()
        for _ in range(self.stack.maxlen):
            self.stack.append(observation)

        return np.array(self.stack).reshape(self.observation_space.low.shape)

    def observation(self, observation):

```

```

        self.stack.append(observation)

        return np.array(self.stack).reshape(self.observation_space.low.shape)

def make_env(env_name, shape=(84,84,1), repeat=4, clip_rewards=False,
             no_ops=0, fire_first=False):
    env = gym.make(env_name)
    env = RepeatActionAndMaxFrame(env, repeat, clip_rewards, no_ops, fire_first)

    env = PreprocessFrame(shape, env)
    env = StackFrames(env, repeat)

    return env

```

```

In [0]: import numpy as np

class ReplayBuffer(object):
    def __init__(self, max_size, input_shape, n_actions):
        self.mem_size = max_size
        self.mem_cntr = 0
        self.state_memory = np.zeros((self.mem_size, *input_shape),
                                     dtype=np.float32)
        self.new_state_memory = np.zeros((self.mem_size, *input_shape),
                                         dtype=np.float32)

        self.action_memory = np.zeros(self.mem_size, dtype=np.int64)
        self.reward_memory = np.zeros(self.mem_size, dtype=np.float32)
        self.terminal_memory = np.zeros(self.mem_size, dtype=np.uint8)

    def store_transition(self, state, action, reward, state_, done):
        index = self.mem_cntr % self.mem_size
        self.state_memory[index] = state
        self.new_state_memory[index] = state_
        self.action_memory[index] = action
        self.reward_memory[index] = reward
        self.terminal_memory[index] = done
        self.mem_cntr += 1

    def sample_buffer(self, batch_size):
        max_mem = min(self.mem_cntr, self.mem_size)
        batch = np.random.choice(max_mem, batch_size, replace=False)

        states = self.state_memory[batch]
        actions = self.action_memory[batch]
        rewards = self.reward_memory[batch]
        states_ = self.new_state_memory[batch]
        terminal = self.terminal_memory[batch]

        return states, actions, rewards, states_, terminal

```

```

In [0]: import tensorflow as tf
        from tensorflow import keras
        import os

class DuelingDeepQNetwork(keras.Model):
    def __init__(self, lr, n_actions, name, input_dims, chkpt_dir):
        super(DuelingDeepQNetwork, self).__init__()
        self.chkpt_dir = chkpt_dir
        self.chkpt_file = os.path.join(self.chkpt_dir, name)
        self.conv1 = keras.layers.Conv2D(input_shape=(-1,*input_dims),
            filters=32, kernel_size=8, data_format='channels_first', strides=4, activation='relu')
        self.conv2 = keras.layers.Conv2D(filters=32, kernel_size=4, strides=2, data_format='channels_first',
            activation='relu')
        self.conv3 = keras.layers.Conv2D(filters=64, kernel_size=3, strides=1, data_format='channels_first',
            activation='relu')
        self.flatlayer = keras.layers.Flatten()
        self.dense1 = keras.layers.Dense(units=1024, activation='relu')
        self.dense2 = keras.layers.Dense(units=512, activation='relu')

        self.V = keras.layers.Dense(units=1, activation='relu')
        self.A = keras.layers.Dense(units=n_actions, activation='relu')
    def call(self, state):
        conv1 = self.conv1(state)
        conv2 = self.conv2(conv1)
        conv3 = self.conv3(conv2)
        flattened = self.flatlayer(conv3)
        dense1 = self.dense1(flattened)
        dense2 = self.dense2(dense1)
        V = self.V(dense2)
        A = self.A(dense2)

        return V, A

    def save_checkpoint(self):
        pass
        print('Saving weights...')
        self.save(self.chkpt_file)

    def load_checkpoint(self):
        pass
        print('loading checkpoint...')
        self.load_weights(self.chkpt_file)

```

```

In [0]: class DuelingDQNAgent(object):
    def __init__(self, gamma, epsilon, lr, n_actions, input_dims,
                 mem_size, batch_size, eps_min=0.01, eps_dec=5e-7,
                 replace=1000, algo=None, env_name=None, chkpt_dir='tmp/dqn'):
        self.gamma = gamma
        self.epsilon = epsilon
        self.lr = lr
        self.n_actions = n_actions
        self.input_dims = input_dims
        self.batch_size = batch_size
        self.eps_min = eps_min
        self.eps_dec = eps_dec
        self.replace_target_cnt = replace
        self.algo = algo
        self.env_name = env_name
        self.chkpt_dir = chkpt_dir
        self.action_space = [i for i in range(n_actions)]
        self.learn_step_counter = 0

        self.memory = ReplayBuffer(mem_size, input_dims, n_actions)

        self.q_eval = DuelingDeepQNetwork(self.lr, self.n_actions,
                                          input_dims=self.input_dims,
                                          name=self.env_name+'_'+self.algo+'_q_eval'
                                          ,
                                          chkpt_dir=self.chkpt_dir)

        self.q_next = DuelingDeepQNetwork(self.lr, self.n_actions,
                                          input_dims=self.input_dims,
                                          name=self.env_name+'_'+self.algo+'_q_next'
                                          ,
                                          chkpt_dir=self.chkpt_dir)

    def choose_action(self, observation):
        if tf.random.uniform([1]) > self.epsilon:
            value, actions = self.q_eval.call(tf.expand_dims(observation, axis=0
            ))

            action = tf.argmax(actions, axis=1)
        else:
            action = np.random.choice(self.action_space)
        return action

    def store_transition(self, state, action, reward, state_, done):
        self.memory.store_transition(state, action, reward, state_, done)

    def sample_memory(self):
        states, actions, rewards, new_states, dones = self.memory.sample_buffer(self.batch_size)
        return states, actions, rewards, new_states, dones

    def replace_target_network(self):
        if self.learn_step_counter & self.replace_target_cnt == 0:
            self.q_next = self.q_eval

    def decrement_epsilon(self):

```

```

self.epsilon = self.epsilon - self.eps_dec \
    if self.epsilon > self.eps_min else self.eps_min

def save_models(self):
    self.q_eval.save_checkpoint()
    self.q_next.save_checkpoint()

def load_models(self):
    self.q_eval.load_checkpoint()
    self.q_next.load_checkpoint()

def learn(self):
    if self.memory.mem_cntr < self.batch_size:
        return

    self.replace_target_network()
    states, actions, rewards, states_, dones = self.sample_memory()
    optimizer = keras.optimizers.RMSprop(learning_rate=self.lr)
    indices = tf.range(self.batch_size)

    with tf.GradientTape() as tape:
        V_s, A_s = self.q_eval.call(states)
        V_s_, A_s_ = self.q_next.call(states_)
        q_pred = tf.add(V_s, (
            A_s - tf.reduce_mean(A_s, axis=1, keepdims=True)
        ))
        q_pred = tf.gather_nd(q_pred, list(zip(indices, actions)))

        q_next = tf.add(V_s_, (
            A_s_ - tf.reduce_mean(A_s_, axis=1, keepdims=True)
        ))
        q_next = tf.reduce_max(q_next, axis=1) * tf.cast(dones, dtype=tf.float32)

        q_target = tf.cast(rewards, dtype=tf.float32) + self.gamma * q_next

        loss = keras.losses.MSE(q_target, q_pred)
        gradient = tape.gradient(loss, self.q_eval.trainable_variables)
        optimizer.apply_gradients(zip(gradient, self.q_eval.trainable_variables))

    self.decrement_epsilon()
    self.learn_step_counter += 1

```

```

In [0]: def main(num_games= 10,load_checkpoint=False, env_name='PongNoFrameskip-v4'):
    env = make_env(env_name)
    best_score = -np.inf

    agent = DuelingDQNAgent(gamma=0.99, epsilon=1.0,lr=0.0001,input_dims=(env.
observation_space.shape),
                            n_actions=env.action_space.n, mem_size=20000, eps_min=0.1
, batch_size=32,replace=100,
                            eps_dec=1e-5, chkpt_dir='models/',algo='DuelingDQNAgent',
env_name=env_name)
    if load_checkpoint:
        agent.load_models()
    fname = agent.algo + '_' + agent.env_name + '_lr' + str(agent.lr) + '_' \
        + str(num_games) + 'games'
    figure_file = 'plots/' + fname + '.png'
    n_steps = 0
    scores, eps_history,steps_array = [], [], []

    for i in range(num_games):
        done = False
        observation = env.reset()
        score = 0
        while not done:
            action = agent.choose_action(observation)
            observation_,reward,done,info = env.step(action)
            score += reward
            if not load_checkpoint:
                agent.store_transition(observation, action, reward,obs
ervation_, int(done))
            agent.learn()

            observation = observation_
            scores.append(score)
            steps_array.append(n_steps)
            avg_score = np.mean(scores[-100:])
            print('episode: ', i,'score: ', score,
                  ' average score %.1f' % avg_score, 'best score %.2f' % best_score
,
                  'epsilon %.2f' % agent.epsilon, 'steps', n_steps)

            if avg_score > best_score:

                best_score = avg_score

            eps_history.append(agent.epsilon)
            if load_checkpoint and n_steps >= 18000:
                break

    x = [i+1 for i in range(len(scores))]
    plot_learning_curve(steps_array, scores, eps_history, figure_file)

```



In [0]: `main(500)`

```
/usr/local/lib/python3.6/dist-packages/gym/logger.py:30: UserWarning: WARN: Box bound precision lowered by casting to float32
```

```
warnings.warn(colorize('%s: %s'%( 'WARN', msg % args), 'yellow'))
```

WARNING:tensorflow:Layer conv2d\_6 is casting an input tensor from dtype float64 to the layer's dtype of float32, which is new behavior in TensorFlow 2. The layer has dtype float32 because it's dtype defaults to floatx.

If you intended to run this layer in float32, you can safely ignore this warning. If in doubt, this warning is likely only an issue if you are porting a TensorFlow 1.X model to TensorFlow 2.

To change all layers to have dtype float64 by default, call `tf.keras.backend.set_floatx('float64')`. To change just this layer, pass `dtype='float64'` to the layer constructor. If you are the author of this layer, you can disable autocasting by passing `autocast=False` to the base Layer constructor.

```
/usr/local/lib/python3.6/dist-packages/gym/envs/atari/atari_env.py:113: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.
```

```
    action = self._action_set[a]
```

```
episode: 0 score: -20.0 average score -20.0 best score -inf epsilon 0.99 s
taps 0
episode: 1 score: -20.0 average score -20.0 best score -20.00 epsilon 0.98
steps 0
episode: 2 score: -21.0 average score -20.3 best score -20.00 epsilon 0.97
steps 0
episode: 3 score: -21.0 average score -20.5 best score -20.00 epsilon 0.97
steps 0
episode: 4 score: -20.0 average score -20.4 best score -20.00 epsilon 0.96
steps 0
episode: 5 score: -20.0 average score -20.3 best score -20.00 epsilon 0.95
steps 0
episode: 6 score: -20.0 average score -20.3 best score -20.00 epsilon 0.94
steps 0
episode: 7 score: -20.0 average score -20.2 best score -20.00 epsilon 0.93
steps 0
episode: 8 score: -20.0 average score -20.2 best score -20.00 epsilon 0.92
steps 0
episode: 9 score: -21.0 average score -20.3 best score -20.00 epsilon 0.91
steps 0
episode: 10 score: -21.0 average score -20.4 best score -20.00 epsilon 0.9
0 steps 0
episode: 11 score: -21.0 average score -20.4 best score -20.00 epsilon 0.8
9 steps 0
episode: 12 score: -20.0 average score -20.4 best score -20.00 epsilon 0.8
8 steps 0
episode: 13 score: -21.0 average score -20.4 best score -20.00 epsilon 0.8
7 steps 0
episode: 14 score: -20.0 average score -20.4 best score -20.00 epsilon 0.8
6 steps 0
episode: 15 score: -21.0 average score -20.4 best score -20.00 epsilon 0.8
5 steps 0
episode: 16 score: -20.0 average score -20.4 best score -20.00 epsilon 0.8
4 steps 0
episode: 17 score: -20.0 average score -20.4 best score -20.00 epsilon 0.8
3 steps 0
episode: 18 score: -21.0 average score -20.4 best score -20.00 epsilon 0.8
3 steps 0
episode: 19 score: -20.0 average score -20.4 best score -20.00 epsilon 0.8
2 steps 0
episode: 20 score: -20.0 average score -20.4 best score -20.00 epsilon 0.8
1 steps 0
episode: 21 score: -21.0 average score -20.4 best score -20.00 epsilon 0.8
0 steps 0
episode: 22 score: -19.0 average score -20.3 best score -20.00 epsilon 0.7
9 steps 0
episode: 23 score: -21.0 average score -20.4 best score -20.00 epsilon 0.7
8 steps 0
episode: 24 score: -20.0 average score -20.4 best score -20.00 epsilon 0.7
7 steps 0
episode: 25 score: -19.0 average score -20.3 best score -20.00 epsilon 0.7
6 steps 0
episode: 26 score: -21.0 average score -20.3 best score -20.00 epsilon 0.7
5 steps 0
episode: 27 score: -21.0 average score -20.4 best score -20.00 epsilon 0.7
4 steps 0
episode: 28 score: -19.0 average score -20.3 best score -20.00 epsilon 0.7
```

```
3 steps 0
episode: 29 score: -20.0 average score -20.3 best score -20.00 epsilon 0.7
2 steps 0
episode: 30 score: -20.0 average score -20.3 best score -20.00 epsilon 0.7
1 steps 0
episode: 31 score: -20.0 average score -20.3 best score -20.00 epsilon 0.7
0 steps 0
episode: 32 score: -21.0 average score -20.3 best score -20.00 epsilon 0.7
0 steps 0
episode: 33 score: -21.0 average score -20.3 best score -20.00 epsilon 0.6
9 steps 0
episode: 34 score: -21.0 average score -20.3 best score -20.00 epsilon 0.6
8 steps 0
episode: 35 score: -21.0 average score -20.4 best score -20.00 epsilon 0.6
7 steps 0
episode: 36 score: -19.0 average score -20.3 best score -20.00 epsilon 0.6
6 steps 0
episode: 37 score: -21.0 average score -20.3 best score -20.00 epsilon 0.6
5 steps 0
episode: 38 score: -21.0 average score -20.4 best score -20.00 epsilon 0.6
4 steps 0
episode: 39 score: -21.0 average score -20.4 best score -20.00 epsilon 0.6
3 steps 0
episode: 40 score: -21.0 average score -20.4 best score -20.00 epsilon 0.6
2 steps 0
episode: 41 score: -21.0 average score -20.4 best score -20.00 epsilon 0.6
1 steps 0
episode: 42 score: -21.0 average score -20.4 best score -20.00 epsilon 0.6
0 steps 0
episode: 43 score: -21.0 average score -20.4 best score -20.00 epsilon 0.5
9 steps 0
episode: 44 score: -19.0 average score -20.4 best score -20.00 epsilon 0.5
8 steps 0
episode: 45 score: -21.0 average score -20.4 best score -20.00 epsilon 0.5
8 steps 0
episode: 46 score: -21.0 average score -20.4 best score -20.00 epsilon 0.5
7 steps 0
episode: 47 score: -21.0 average score -20.4 best score -20.00 epsilon 0.5
6 steps 0
episode: 48 score: -20.0 average score -20.4 best score -20.00 epsilon 0.5
5 steps 0
episode: 49 score: -20.0 average score -20.4 best score -20.00 epsilon 0.5
4 steps 0
episode: 50 score: -21.0 average score -20.4 best score -20.00 epsilon 0.5
3 steps 0
episode: 51 score: -21.0 average score -20.4 best score -20.00 epsilon 0.5
2 steps 0
episode: 52 score: -20.0 average score -20.4 best score -20.00 epsilon 0.5
1 steps 0
episode: 53 score: -21.0 average score -20.4 best score -20.00 epsilon 0.5
0 steps 0
episode: 54 score: -21.0 average score -20.5 best score -20.00 epsilon 0.4
9 steps 0
episode: 55 score: -20.0 average score -20.4 best score -20.00 epsilon 0.4
8 steps 0
episode: 56 score: -21.0 average score -20.5 best score -20.00 epsilon 0.4
7 steps 0
```

```
episode: 57 score: -21.0 average score -20.5 best score -20.00 epsilon 0.4
6 steps 0
episode: 58 score: -21.0 average score -20.5 best score -20.00 epsilon 0.4
5 steps 0
episode: 59 score: -21.0 average score -20.5 best score -20.00 epsilon 0.4
5 steps 0
episode: 60 score: -19.0 average score -20.5 best score -20.00 epsilon 0.4
3 steps 0
episode: 61 score: -21.0 average score -20.5 best score -20.00 epsilon 0.4
3 steps 0
episode: 62 score: -21.0 average score -20.5 best score -20.00 epsilon 0.4
1 steps 0
episode: 63 score: -21.0 average score -20.5 best score -20.00 epsilon 0.4
1 steps 0
episode: 64 score: -21.0 average score -20.5 best score -20.00 epsilon 0.4
0 steps 0
episode: 65 score: -21.0 average score -20.5 best score -20.00 epsilon 0.3
9 steps 0
episode: 66 score: -21.0 average score -20.5 best score -20.00 epsilon 0.3
8 steps 0
episode: 67 score: -21.0 average score -20.5 best score -20.00 epsilon 0.3
7 steps 0
episode: 68 score: -20.0 average score -20.5 best score -20.00 epsilon 0.3
6 steps 0
episode: 69 score: -20.0 average score -20.5 best score -20.00 epsilon 0.3
5 steps 0
episode: 70 score: -21.0 average score -20.5 best score -20.00 epsilon 0.3
4 steps 0
episode: 71 score: -21.0 average score -20.5 best score -20.00 epsilon 0.3
3 steps 0
episode: 72 score: -21.0 average score -20.5 best score -20.00 epsilon 0.3
2 steps 0
episode: 73 score: -21.0 average score -20.5 best score -20.00 epsilon 0.3
1 steps 0
episode: 74 score: -21.0 average score -20.5 best score -20.00 epsilon 0.3
1 steps 0
episode: 75 score: -21.0 average score -20.5 best score -20.00 epsilon 0.3
0 steps 0
episode: 76 score: -21.0 average score -20.5 best score -20.00 epsilon 0.2
9 steps 0
episode: 77 score: -21.0 average score -20.6 best score -20.00 epsilon 0.2
8 steps 0
episode: 78 score: -21.0 average score -20.6 best score -20.00 epsilon 0.2
7 steps 0
episode: 79 score: -21.0 average score -20.6 best score -20.00 epsilon 0.2
6 steps 0
episode: 80 score: -21.0 average score -20.6 best score -20.00 epsilon 0.2
6 steps 0
episode: 81 score: -21.0 average score -20.6 best score -20.00 epsilon 0.2
5 steps 0
episode: 82 score: -21.0 average score -20.6 best score -20.00 epsilon 0.2
4 steps 0
episode: 83 score: -21.0 average score -20.6 best score -20.00 epsilon 0.2
3 steps 0
episode: 84 score: -21.0 average score -20.6 best score -20.00 epsilon 0.2
2 steps 0
episode: 85 score: -21.0 average score -20.6 best score -20.00 epsilon 0.2
```

```
1 steps 0
episode: 86 score: -21.0 average score -20.6 best score -20.00 epsilon 0.2
0 steps 0
episode: 87 score: -21.0 average score -20.6 best score -20.00 epsilon 0.1
9 steps 0
episode: 88 score: -21.0 average score -20.6 best score -20.00 epsilon 0.1
8 steps 0
```